

Agentic AI | ISP


December 3, 2025

Executive Summary

- Most enterprises have adopted GenAI tools, yet few have realized meaningful ROI because copilots remain shallowly integrated into workflows and lack the trust, memory, and guardrails needed for real operational impact
- This gap motivated our project: to understand where *agentic* AI, systems that perceive, plan, act, and iterate inside workflows, could unlock true business value
- Our early work in search-funds seemed promising at first: large task volume, clear success metrics, and heavy reliance on research and personalization. But as we built early prototypes, it became clear that increased lead volume did not translate to ROI
- We pivoted to consulting, where workflow structure and economics align more directly: efficiency gains translate into real value through additional case capacity and reduced analyst burden
- Across management consulting workflows, we found multiple high-volume, repeatable tasks well-suited for agentic automation, from primary research to secondary analysis, surveys, and slide creation
- We built an MVP focused on primary research automation: the highest-signal, highest-volume input in consulting
- From this work, four key lessons emerged:
 - **1. Integration friction, not model capability, is the true bottleneck**
 - **2. Platform giants will commoditize generic actions; defensibility lies in workflow expertise**
 - **3. Market selection is critical: costly integrations require large, repeatable verticals**
 - **4. Trust, not speed, is the gating factor for enterprise adoption**
- Together, these insights outline a clear path forward: agentic AI can transform knowledge-work industries, but only when paired with careful market selection and deep workflow integration

The Problem

GenAI adoption has delivered limited ROI impact due to poor workflow integration and low trust, pointing to a major opportunity for more capable, workflow-aware systems




Today: Horizontal Use Cases = Adoption without ROI

- Focused on **generic copilots and chatbots** (e.g., ChatGPT, Copilot) with diffuse benefits.
- Deliver **incremental boosts to individual productivity**, but little measurable P&L impact.
- Struggle in **enterprise workflows** due to lack of memory, integration, and workflow alignment.
- Leave companies **trapped in the “GenAI Divide”**, with adoption but no transformation.

Future: Vertical Use Cases = ROI through Workflow Fit


- Designed for **function-specific workflows**, not just general productivity.
- Powered by **persistent memory and domain context** to handle complex, multi-step tasks.
- Validated by **SMEs embedded in day-to-day processes**, ensuring accuracy and adoption.
- Integrated into **legacy systems and data pipelines** to drive scale and resilience.
- Unlocks **P&L impact** by linking AI output to business outcomes, poised to overcome today’s “GenAI Divide.”

Agentic AI emerged as the natural next step to bridge these gaps: systems designed to perceive, plan, act, & iterate within human workflows



Our Journey

Search-fund sourcing initially looked ideal for AI agents: large volume of repetitive tasks requiring reasoning with clear success metrics



However, while agents could generate more leads, lead volume isn't the real bottleneck for searchers; ROI would be low, making this space unattractive for us.


By contrast, consulting firms can turn efficiency gains from agentic AI into real ROI by taking on additional clients and cases firms previously lacked capacity to serve

ROI from Agentic AI adoption in Consulting	Revenue Generation	Notes:
	Upsell existing clients	<< Don't believe that we can upsell existing clients further
	Gain new clients	<< Additional volume; could offer a lower priced solution that top tier consulting firms can use brand recognition to compete on
	Cost Savings	
	Headcount reduction	<< Consultants are unlikely to let go of people unless they can no longer get more revenue
	Time spent per output of work created	

Consulting is a domain where several core workflows are both repeatable and high-volume, making them naturally suited for agentic automation


Due Diligence Workstreams	Example workflow	Agentic AI disruption potential	Prioritization
Primary Research	Slide update based on call notes	5/5	1 - most time spent, easiest to implement, large disruption potential
Secondary Research	TAM / SAM model	3.5/5	2 - significant time spent, numerous complexities to execute correctly, ripe for learning and impact
Survey	Testing survey design / logic	4.5/5	3 - subset of primary; better to tackle holistically
Deck / Storylining	Blank slide / hypothesis generation	3/5	4 - significant judgement required
Case Management	Internal workload management	1/5	5 - people-centric, less suitable for agent

Our project focuses on primary research automation: the highest-signal, highest-volume input in consulting workflows




Our MVP


Our architecture decomposes consulting tasks into distinct agent roles



Our MVP connects each agent via an orchestrated chain to complete tasks up to consultant standards




n8n Deep-Dive: To operationalize this architecture, we built a detailed n8n workflow that makes each agent's responsibilities explicit and exposes the technical interactions under the hood



We define clear success criteria and will evaluate the system on both technical fidelity and the real value it delivers to consulting teams

Technical Verification (Does the System Function?)	User Validation (Does the Consulting Team Benefit?)
Focus: The Agentic Loop & Architecture	Focus: The Management Consulting Workflow
Metric 1: Tool Selection Accuracy Did the agent pick the right tool for the prompt?	Metric 1: Creation Velocity Reduction in time required to produce a 'First Draft' deck.
Metric 2: RAG Groundedness Is the data retrieved by n8n factually accurate to the source?	Metric 2: The "Undo" Rate How often does the user reject or delete the agent's output?
Metric 3: VSTO Execution Success Did the slide update without syntax errors or crashes?	Metric 3: Learning Curve Can a user create complex slides without learning specific syntax?
Actions: Deterministic unit tests for API payloads/responses; Deterministic integration tests for end to end actions; LLM-as-a-Judge tests to grade reasoning traces and output quality.	Actions: Measure use of product by management consulting teams; Blind quality audits.

Here's a sneak peek at the system in action...



We encountered three technical bottlenecks in our MVP design; for each we defined a clear path for success and outlined next steps

Bottleneck	PowerPoint API Not Agent-Friendly	Version control	n8n speed bottleneck
Description	PowerPoint's native interface is too low-level for an agent to use directly, and it provides almost no guardrails. Even simple edits require detailed, fragile instructions that agents can easily get wrong.	PowerPoint has no native diffing or version visualization , so there's no clean way to track changes the way tools like Cursor or Git do	n8n becomes a bottleneck for complex multi-agent workflows, especially as the number of nodes grows and workflows become state-heavy.
Solution	We created a library of simple , predefined slide-editing actions (e.g., add_bullet_point) that add structure and guardrails. The agent can now call safe, reliable actions instead of working with low-level PowerPoint logic.	We built a two-window diff view that highlights new changes on the updated deck and includes an approve / reject interaction	Shift orchestration and agent execution to LangChain , which offers faster routing, parallelization, and more robust agent-tooling for long workflows.
Next steps	Expand the internal library of custom methods , expose them to the slide-updater agent, and run performance/accuracy tests to ensure the agent still behaves predictably at scale.	Define and standardize the visual language of “highlights” for different content types (charts, tables, images, text boxes) and extend the diff logic to more complex slide elements.	Define the final agent architecture inside LangChain, benchmark latency across typical tasks, and iteratively migrate the full pipeline once performance targets are met.

Key Insights

Key Insight #1: Agentic AI is powerful for consulting workflows, but integration friction is the real bottleneck

- Agents performed extremely well on extraction, synthesis and quantitative updates, especially for transcript-driven workflows
- The biggest blockers weren't model capability but enterprise integration details:
 - PowerPoint APIs are not agent friendly → required writing custom methods
 - No version control, no accept/reject flow → consultants cannot trust outputs without structured diffs
 - n8n orchestration proved slow → memoryless workflows limited reliability
- **Insight:** The “last mile” (PowerPoint, speed, versioning, human-in-the-loop (HITL) review) determines adoption more than the AI itself

Key Insight #2: Platform giants (e.g., MSFT) will own generic integrations; moat lies in deep workflow expertise

- Office products already have privileged write-access, data permissions, user identity, and native slide rendering; our PowerPoint work illustrated how hard this is to replicate externally
- Copilot will likely commoditize actions like note-taking, slide drafting, basic extraction
- Our defensibility comes from:
 - Knowledge of actual consulting workflows, not generic productivity
 - Granularity of tasks (e.g., call-extract → hypothesis generation → storyline integration) that are invisible to Microsoft
 - High-judgment patterns and domain-specific agent chaining
- **Insight:** The moat is verticalization; building intellectual property around how analysts really work, not horizontal features

Key Insight #3: Market selection matters: costly, complex integration only pays off when selling into large, repeatable verticals

- Consulting is an attractive workflow fit but a difficult commercial market:
 - Limited number of major buyers
 - Slow procurement
 - High security requirements
- Because building integrations is complex (APIs, HITL, versioning), success requires markets with:
 - Many buyers
 - Standardized workflows
- **Insight:** If the integration tax is high, we must spread that tax across a large number of customers, hence the need to focus on scalable industries

Key Insight #4: Trust, not speed, is the gating factor for enterprise adoption

- Across the research and literature review:
 - Enterprises need explainability, traceability, and rollback paths
 - High-risk edits (facts, numbers, quotes) require HITL review; low-risk changes can be automated
- Our MVP showed the path:
 - Enables an accept/edit/reject flow similar to Git + Cursor inline edits
 - Creates a supervised workflow: humans communicate intent, agents draft, humans review
- **Insight:** Adoption only happens when customers feel in control; trust infrastructure matters at least as much as model intelligence

Open Questions

Economic & Workflow

- *What magnitude of productivity gain (2x, 5x, 10x) is necessary to justify switching costs and overcome entrenched legacy workflows?*
- *At what point does the cognitive load of auditing an agent's output exceed the time savings from automation?*
- *How do firms redesign workflows when automation alters the pace, sequencing, and unit economics of consulting work?*
- *What new commercial models (e.g., fixed-fee, outcome-based) emerge as the marginal cost of producing deliverables approaches zero?*

Organization & Talent

- *If AI automates apprenticeship tasks, how do firms develop future leaders and judgment?*
- *Where should mid-level talent be redeployed when slide production disappears — client time or more workstreams?*
- *How should firms balance headcount reduction with the need to maintain capability-building pipelines?*
- *As workflows evolve, what skills become most critical for consultants to grow into senior roles?*

Go-to-Market

- *What does vertical defensibility look like as platform giants rapidly expand their native capabilities?*
- *Which large, repeatable verticals justify the high integration tax?*
- *How do we identify markets with enough long-tail demand to scale integration costs?*